

Heart Disease (including Coronary Heart Disease, Hypertension, and Stroke) remains the No. 1 cause of death in the US. The Heart Disease and Stroke Statistics—2019 Update from the American Heart Association indicates that:

In this machine learning project, we have collected the dataset from UCI ([https://archive.ics.uci.edu/ml/datasets/statlog+\(heart\)](https://archive.ics.uci.edu/ml/datasets/statlog+(heart))) and we will be using Machine Learning to make predictions on whether a person is suffering from Heart Disease or not.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage.
# It includes your credentials.
# You might want to remove those credentials before you share the
# notebook.
cos_client = ibm_boto3.client(service_name='s3',
                              ibm_api_key_id='5i6ow0P4tpnCXRwPmHgJ0lvS_22JMaJwUe39D0IQ-wz8',
                              ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
                              config=Config(signature_version='oauth'),
                              endpoint_url='https://s3.private.us.cloud-object-
storage.appdomain.cloud')

bucket = 'visualizingandpredictingheartdise-donotdelete-pr-
dhoozkeqjwqosc'
object_key = 'Heart_Disease_Prediction.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)['Body']
# add missing __iter__ method, so pandas accepts body as file-like
object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType(
__iter__, body )

data = pd.read_csv(body)
data.head()
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG
0	70	1	4	130	322	0	
2							
1	67	0	3	115	564	0	
2							
2	57	1	2	124	261	0	
0							
3	64	1	4	128	263	0	
0							
4	74	0	2	120	269	0	
2							

	Max HR	Exercise angina	ST depression	Slope of ST	\
0	109	0	2.4	2	
1	160	0	1.6	2	
2	141	0	0.3	1	
3	105	1	0.2	2	
4	121	1	0.2	1	

	Number of vessels fluro	Thallium	Heart Disease
0	3	3	Presence
1	0	7	Absence
2	0	7	Presence
3	1	7	Absence
4	1	3	Absence

data.tail()

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG
265	52	1	3	172	199	1	
0							
266	44	1	2	120	263	0	
0							
267	56	0	2	140	294	0	
2							
268	57	1	4	140	192	0	
0							
269	67	1	4	160	286	0	
2							

	Max HR	Exercise angina	ST depression	Slope of ST	\
265	162	0	0.5	1	
266	173	0	0.0	1	
267	153	0	1.3	2	
268	148	0	0.4	2	
269	108	1	1.5	2	

	Number of vessels fluro	Thallium	Heart Disease
265	3	3	Presence
266	0	7	Absence
267	0	7	Presence
268	1	7	Absence
269	1	3	Absence

265	0	7	Absence
266	0	7	Absence
267	0	3	Absence
268	0	6	Absence
269	3	3	Presence

```
print(f'No of Rows in the dataset : {data.shape[0]}')
print(f'No of Columns in the dataset : {data.shape[1]}')
```

```
No of Rows in the dataset : 270
No of Columns in the dataset : 14
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Age                                  270 non-null    int64
 1   Sex                                  270 non-null    int64
 2   Chest pain type                      270 non-null    int64
 3   BP                                    270 non-null    int64
 4   Cholesterol                          270 non-null    int64
 5   FBS over 120                        270 non-null    int64
 6   EKG results                         270 non-null    int64
 7   Max HR                              270 non-null    int64
 8   Exercise angina                     270 non-null    int64
 9   ST depression                       270 non-null    float64
10   Slope of ST                         270 non-null    int64
11   Number of vessels fluro             270 non-null    int64
12   Thallium                            270 non-null    int64
13   Heart Disease                       270 non-null    object
dtypes: float64(1), int64(12), object(1)
memory usage: 29.7+ KB
```

```
data.describe()
```

	Age	Sex	Chest pain type	BP
Cholesterol \				
count	270.000000	270.000000	270.000000	270.000000
mean	54.433333	0.677778	3.174074	131.344444
std	9.109067	0.468195	0.950090	17.861608
min	29.000000	0.000000	1.000000	94.000000
25%	48.000000	0.000000	3.000000	120.000000
50%	55.000000	1.000000	3.000000	130.000000

245.000000				
75%	61.000000	1.000000	4.000000	140.000000
280.000000				
max	77.000000	1.000000	4.000000	200.000000
564.000000				


	FBS over 120	EKG results	Max HR	Exercise angina	ST
depression \					
count	270.000000	270.000000	270.000000	270.000000	270.000000
mean	0.148148	1.022222	149.677778	0.329630	1.05000
std	0.355906	0.997891	23.165717	0.470952	1.14521
min	0.000000	0.000000	71.000000	0.000000	0.00000
25%	0.000000	0.000000	133.000000	0.000000	0.00000
50%	0.000000	2.000000	153.500000	0.000000	0.80000
75%	0.000000	2.000000	166.000000	1.000000	1.60000
max	1.000000	2.000000	202.000000	1.000000	6.20000

	Slope of ST	Number of vessels fluro	Thallium
count	270.000000	270.000000	270.000000
mean	1.585185	0.670370	4.696296
std	0.614390	0.943896	1.940659
min	1.000000	0.000000	3.000000
25%	1.000000	0.000000	3.000000
50%	2.000000	0.000000	3.000000
75%	2.000000	1.000000	7.000000
max	3.000000	3.000000	7.000000


data.columns

```
Index(['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120',
      'EKG results', 'Max HR', 'Exercise angina', 'ST depression',
      'Slope of ST', 'Number of vessels fluro', 'Thallium', 'Heart Disease'],
      dtype='object')
```

```
sns.displot(x=data["Age"], kde=True, color='green')
plt.title("AGE distribution");
```



```
labels = ['Male', 'Female']  
order = data['Sex'].value_counts().index  
  
sns.countplot(x='Sex', data=data, order=order)  
plt.xticks([0, 1], labels)  
plt.show()
```



```
data['Chest pain type'].value_counts()
```

```
4    129
3     79
2     42
1     20
```

```
Name: Chest pain type, dtype: int64
```


```
labels = ["asymptomatic", "non-anginal pain", "atypical angina", "typical  
angina"]
```

```
order = data['Chest pain type'].value_counts().index
```


```
sns.countplot(x='Chest pain type', data=data, order=order)
```

```
plt.xticks([0,1,2,3], labels, rotation=45)
```


```
plt.show()
```



```
sns.histplot(x=data["BP"], kde=True)  
plt.title("BP");
```



```
sns.histplot(x=data["Cholesterol"],kde=True)
plt.title("Cholesterol");
```

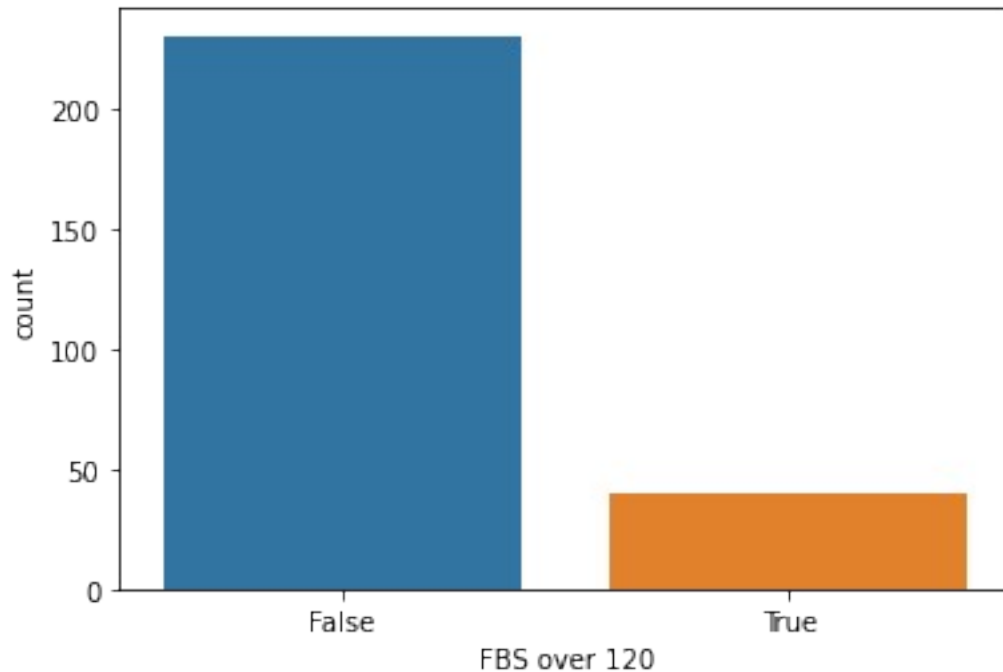


```
data['FBS over 120'].value_counts()

0    230
1     40
Name: FBS over 120, dtype: int64

labels = ["False", 'True']
order = data['FBS over 120'].value_counts().index

sns.countplot(x='FBS over 120', data=data, order=order)
plt.xticks([0,1], labels=labels)
plt.show()
```

```
data['EKG results'].value_counts()
```

```
2    137
```

```
0    131
```

```
1      2
```

```
Name: EKG results, dtype: int64
```

```
labels = ["showing probable or definite left ventricular hypertrophy  
by Estes' criteria", "normal",  
          "having ST-T wave abnormality"]
```


```
order = data['EKG results'].value_counts().index
```

```
plt.figure(figsize=(10, 8))
```


```
sns.countplot(x='EKG results', data=data, order=order)
```

```
#plt.xticks([0,1,2], labels=labels)
```

```
plt.show()
```



```
sns.histplot(x=data["Max HR"],kde=True)  
plt.title("Maximum Heart Rate");
```



```
data['Exercise angina'].value_counts()
```

```
0    181
```

```
1     89
```

```
Name: Exercise angina, dtype: int64
```

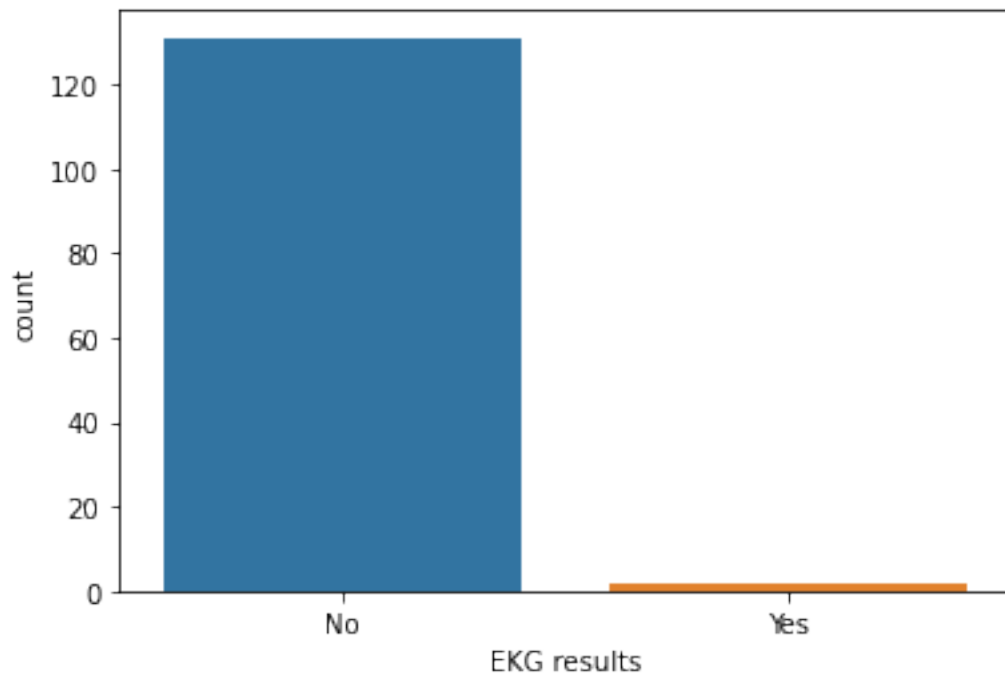
```
labels = ["No", "Yes"]
```

```
order = data['Exercise angina'].value_counts().index
```


```
sns.countplot(x='EKG results', data=data, order=order)
```

```
plt.xticks([0,1], labels=labels)
```

```
plt.show()
```



```
sns.histplot(x=data["ST depression"],kde=True)  
plt.title("ST depression");
```



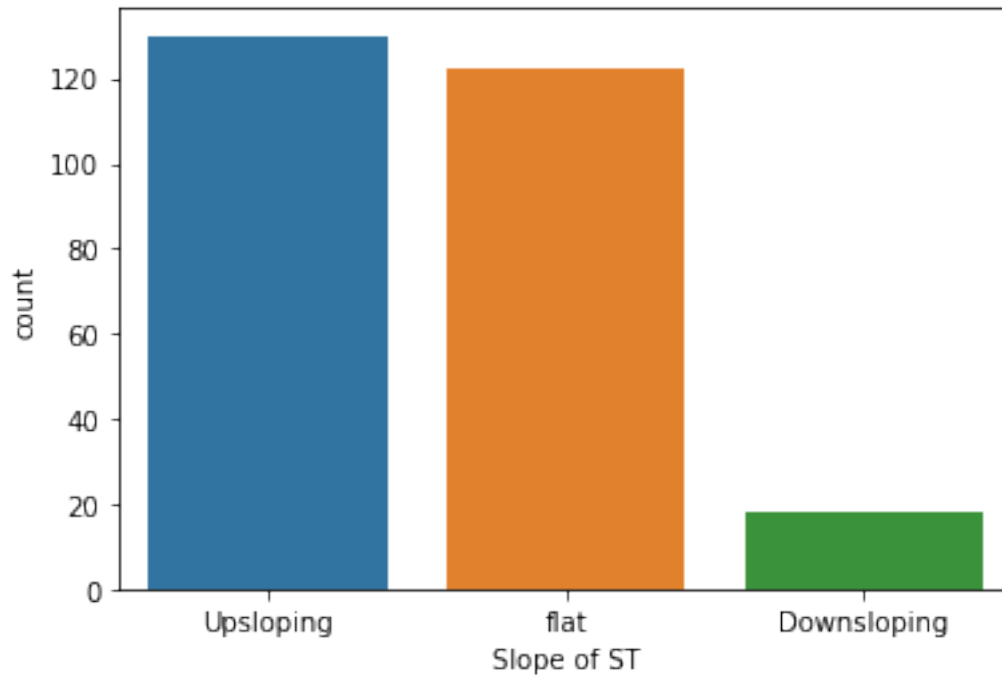
```
data['Slope of ST'].value_counts()
```

```
1    130  
2    122
```

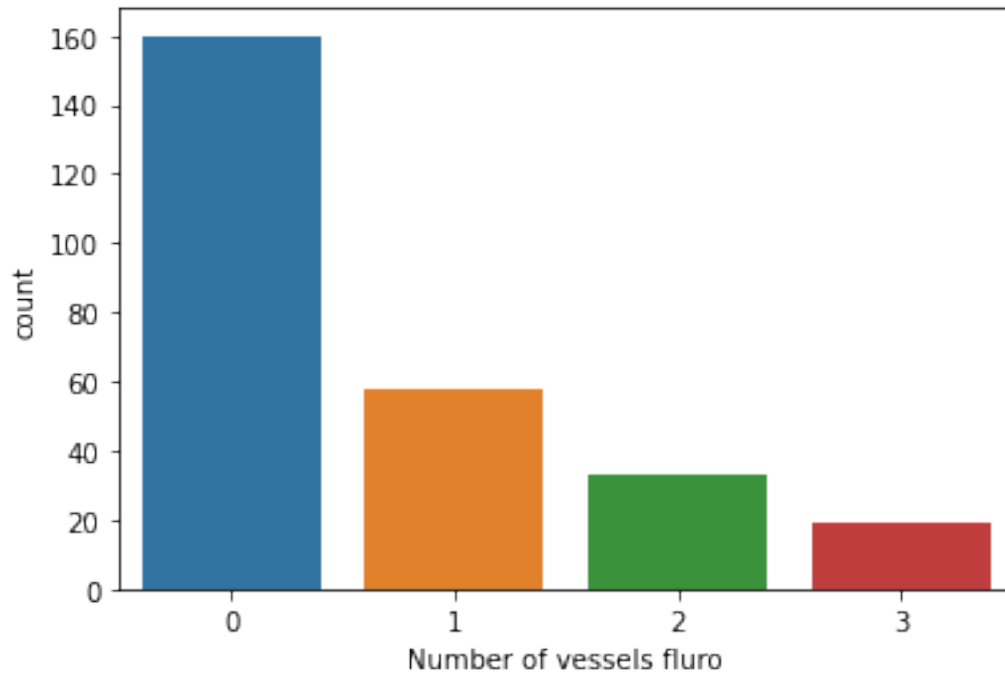
```
3      18
Name: Slope of ST, dtype: int64

labels = ["Upsloping", "flat", "Downsloping"]
order = data['Slope of ST'].value_counts().index

sns.countplot(x='Slope of ST', data=data, order=order)
plt.xticks([0,1,2], labels=labels)
plt.show()
```



```
sns.countplot(x='Number of vessels fluoro', data=data);
```



```
data['Thallium'].value_counts()
```

```
3    152
```

```
7    104
```

```
6     14
```

```
Name: Thallium, dtype: int64
```

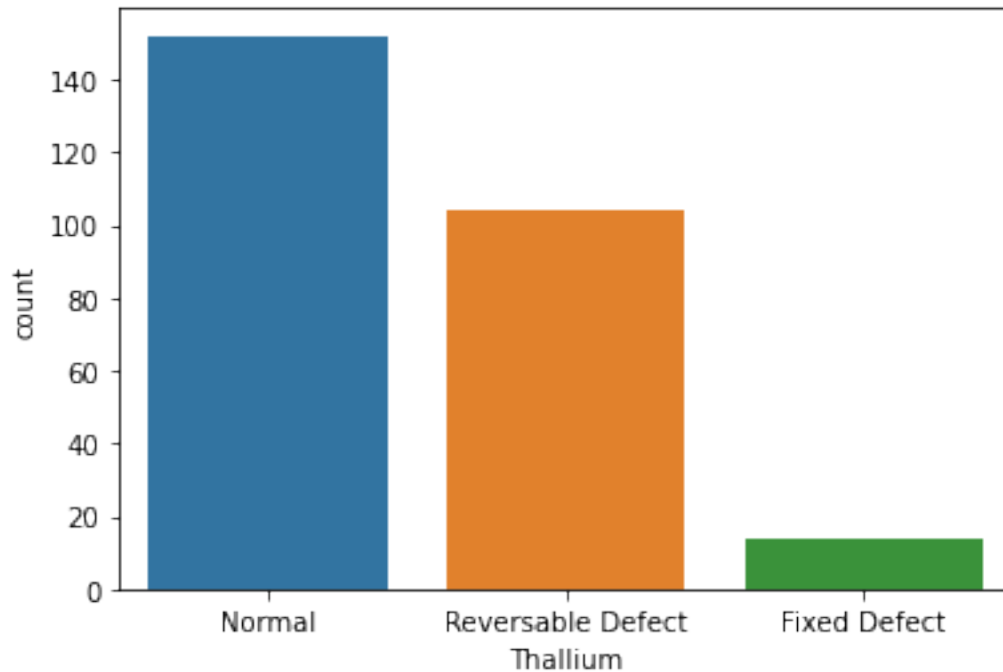
```
labels = ["Normal", "Reversible Defect", "Fixed Defect"]
```

```
order = data['Thallium'].value_counts().index
```

```
sns.countplot(x='Thallium', data=data, order=order)
```

```
plt.xticks([0,1,2], labels=labels)
```

```
plt.show()
```

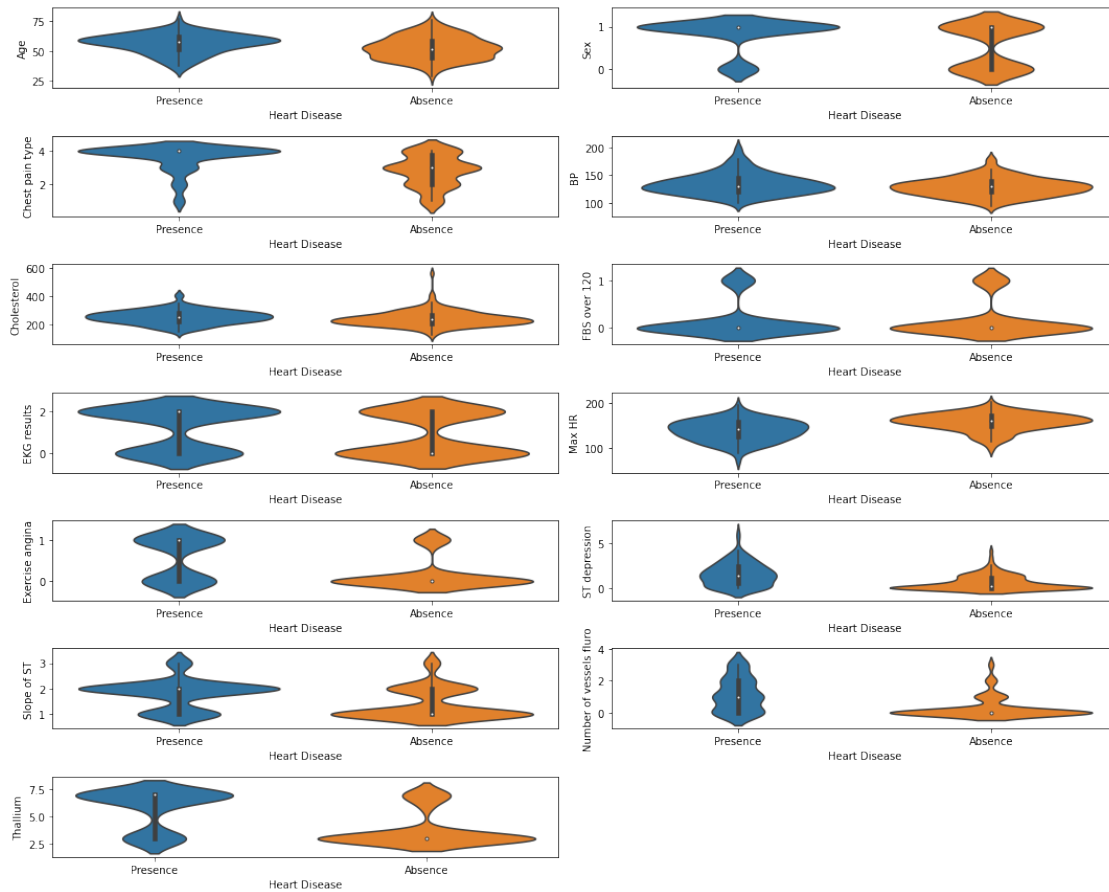


```
cols = ['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS
over 120',
        'EKG results', 'Max HR', 'Exercise angina', 'ST depression',
        'Slope of ST', 'Number of vessels fluro', 'Thallium']
cols
```

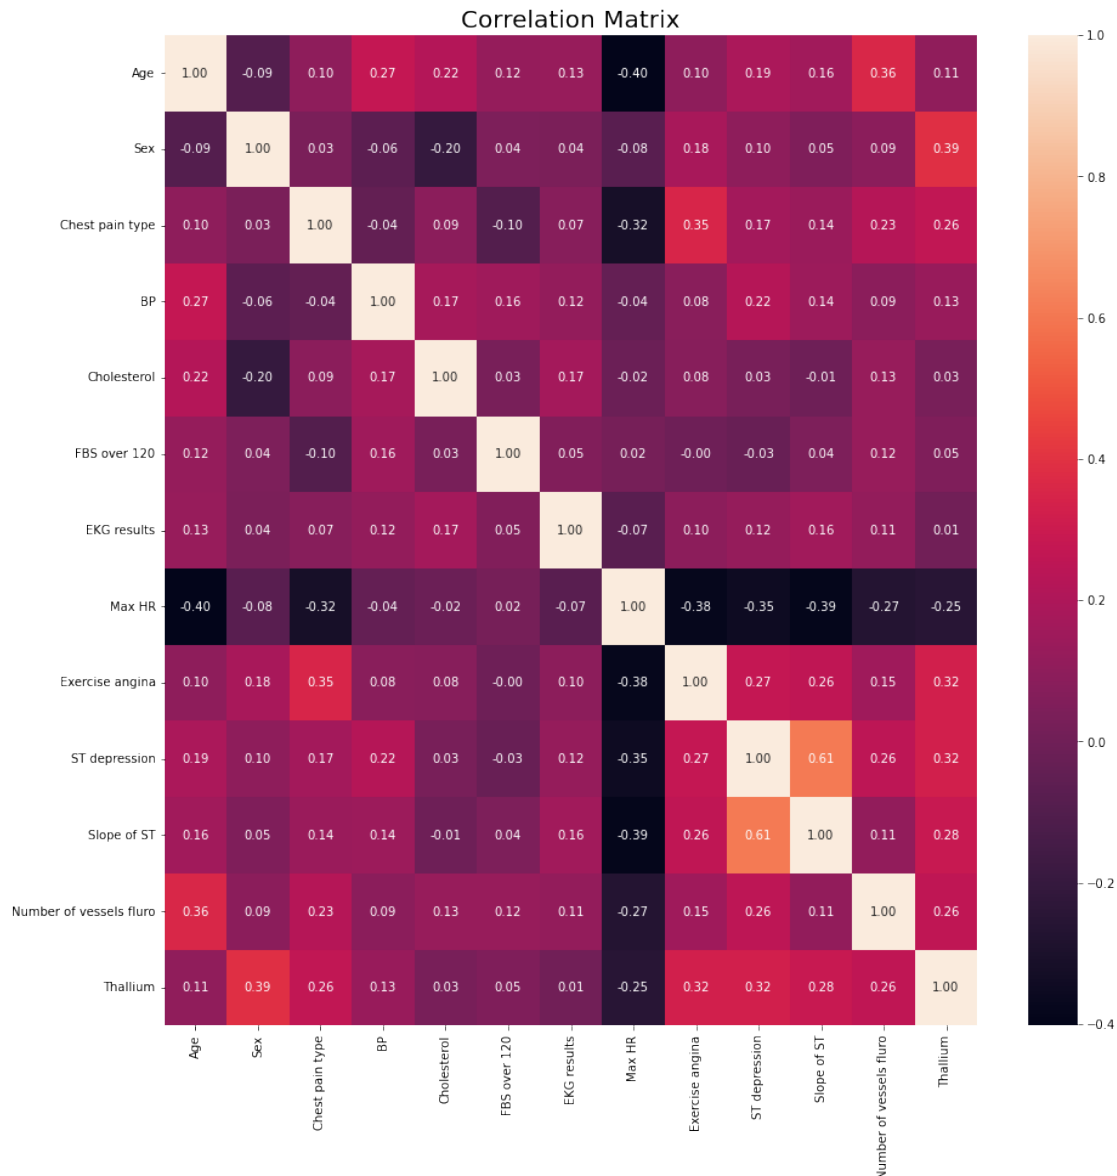
```
['Age',
 'Sex',
 'Chest pain type',
 'BP',
 'Cholesterol',
 'FBS over 120',
 'EKG results',
 'Max HR',
 'Exercise angina',
 'ST depression',
 'Slope of ST',
 'Number of vessels fluro',
 'Thallium']
```

```
plt.figure(figsize=(15, 12))
for i in range(len(cols)):
    plt.subplot(7, 2, i+1)
    sns.violinplot(x="Heart Disease", y=cols[i], data=data)
plt.tight_layout()
plt.plot()
```

```
[]
```



```
corr_matrix = data.corr()
plt.figure(figsize = (15, 15))
sns.heatmap(corr_matrix,annot=True,fmt='0.2f')
plt.title("Correlation Matrix", fontsize = 20)
plt.show()
```

```
data.columns
```

```
Index(['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120',
      'EKG results', 'Max HR', 'Exercise angina', 'ST depression',
      'Slope of ST', 'Number of vessels fluro', 'Thallium', 'Heart Disease'],
      dtype='object')
```

```
x = data.iloc[:,0:13]
```

```
y = data['Heart Disease'].replace({"Absence":0,"Presence":1})
```

```
x.head()
```

```
Age Sex Chest pain type BP Cholesterol FBS over 120 EKG
results \
```

0	70	1	4	130	322	0
2						
1	67	0	3	115	564	0
2						
2	57	1	2	124	261	0
0						
3	64	1	4	128	263	0
0						
4	74	0	2	120	269	0
2						

	Max HR	Exercise angina	ST depression	Slope of ST	\
0	109	0	2.4	2	
1	160	0	1.6	2	
2	141	0	0.3	1	
3	105	1	0.2	2	
4	121	1	0.2	1	

	Number of vessels	fluro	Thallium
0		3	3
1		0	7
2		0	7
3		1	7
4		1	3

y

0	1
1	0
2	1
3	0
4	0

	..
265	0
266	0
267	0
268	0
269	1

Name: Heart Disease, Length: 270, dtype: int64

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,
test_size=0.10)
```

```
print(f'x_train contains: {x_train.shape[0]} rows and
{x_train.shape[1]} columns')
print(f'x_test contains: {x_test.shape[0]} rows and {x_test.shape[1]}
columns')
```

x_train contains: 243 rows and 13 columns

x_test contains: 27 rows and 13 columns

```

print(f'y_train contains: {y_train.shape}')
print(f'y_test contains: {y_test.shape}')

y_train contains: (243,)
y_test contains: (27,)

print(f'x_train contains: {x_train.shape[0]} rows and
{x_train.shape[1]} columns')
print(f'x_test contains: {x_test.shape[0]} rows and {x_test.shape[1]}
columns')

x_train contains: 243 rows and 13 columns
x_test contains: 27 rows and 13 columns

print(f'y_train contains: {y_train.shape}')
print(f'y_test contains: {y_test.shape}')

y_train contains: (243,)
y_test contains: (27,)

from sklearn.metrics import classification_report
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier()
knn.fit(x_train,y_train)

KNeighborsClassifier()

knn.score(x_test,y_test)

0.5185185185185185

knn_pred = knn.predict(x_test)

len(knn_pred)

27

print(classification_report(y_test,knn_pred))

              precision    recall  f1-score   support

     0       0.71      0.53      0.61         19
     1       0.31      0.50      0.38          8

 accuracy          0.52         27
 macro avg       0.51      0.51      0.49         27
weighted avg       0.59      0.52      0.54         27

print(classification_report(y_train, knn.predict(x_train)))

```

	precision	recall	f1-score	support
0	0.77	0.85	0.81	131
1	0.80	0.71	0.75	112
accuracy			0.78	243
macro avg	0.78	0.78	0.78	243
weighted avg	0.78	0.78	0.78	243

```
x_test.shape
```

```
(27, 13)
```

```
from ibm_watson_machine_learning import APIClient
wml_credentials={
    "url": "https://us-south.ml.cloud.ibm.com",
    "apikey": "X2Pk1QhBMW2t1ZpNhVtmi3RShe8Icb-61H84qlBy3UyG"
}
client=APIClient(wml_credentials)

def guid_from_space_name(client, space_name):
    space = client.spaces.get_details()
    #print(space)
    return(next(item for item in space['resources'] if item['entity']
["name"] == space_name)['metadata']['id'])

space_uid = guid_from_space_name(client, 'models')
print("Space UID = " + space_uid)

Space UID = e74734cd-b84a-4c71-90d2-ceefc839e963

client.set.default_space(space_uid)

'SUCCESS'

client.software_specifications.list()
```

```
-----
----
NAME                               ASSET_ID
TYPE
default_py3.6                     0062b8c9-8b7d-44a0-a9b9-46c416adcbd9
base
kernel-spark3.2-scala2.12         020d69ce-7ac1-5e68-ac1a-31189867356a
base
pytorch-onnx_1.3-py3.7-edt        069ea134-3346-5748-b513-49120e15d288
base
scikit-learn_0.20-py3.6           09c5a1d0-9c1e-4473-a344-eb7b665ff687
base
spark-mllib_3.0-scala_2.12        09f4cff0-90a7-5899-b9ed-1ef348aebdee
base
```

pytorch-onnx_rt22.1-py3.9	0b848dd4-e681-5599-be41-b5f6fccc6471
base	
ai-function_0.1-py3.6	0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda
base	
shiny-r3.6	0e6e79df-875e-4f24-8ae9-62dcc2148306
base	
tensorflow_2.4-py3.7-horovod	1092590a-307d-563d-9b62-4eb7d64b3f22
base	
pytorch_1.1-py3.6	10ac12d6-6b30-4ccd-8392-3e922c096a92
base	
tensorflow_1.15-py3.6-ddl	111e41b3-de2d-5422-a4d6-bf776828c4b7
base	
autoai-kb_rt22.2-py3.10	125b6d9a-5b1f-5e8d-972a-b251688ccf40
base	
runtime-22.1-py3.9	12b83a17-24d8-5082-900f-0ab31fbfd3cb
base	
scikit-learn_0.22-py3.6	154010fa-5b3b-4ac1-82af-4d5ee5abbc85
base	
default_r3.6	1b70aec3-ab34-4b87-8aa0-a4a3c8296a36
base	
pytorch-onnx_1.3-py3.6	1bc6029a-cc97-56da-b8e0-39c3880dbbe7
base	
kernel-spark3.3-r3.6	1c9e5454-f216-59dd-a20e-474a5cdf5988
base	
pytorch-onnx_rt22.1-py3.9-edt	1d362186-7ad5-5b59-8b6c-9d0880bde37f
base	
tensorflow_2.1-py3.6	1eb25b84-d6ed-5dde-b6a5-3fbdf1665666
base	
spark-mllib_3.2	20047f72-0a98-58c7-9ff5-a77b012eb8f5
base	
tensorflow_2.4-py3.8-horovod	217c16f6-178f-56bf-824a-b19f20564c49
base	
runtime-22.1-py3.9-cuda	26215f05-08c3-5a41-a1b0-da66306ce658
base	
do_py3.8	295addb5-9ef9-547e-9bf4-92ae3563e720
base	
autoai-ts_3.8-py3.8	2aa0c932-798f-5ae9-abd6-15e0c2402fb5
base	
tensorflow_1.15-py3.6	2b73a275-7cbf-420b-a912-eae7f436e0bc
base	
kernel-spark3.3-py3.9	2b7961e2-e3b1-5a8c-a491-482c8368839a
base	
pytorch_1.2-py3.6	2c8ef57d-2687-4b7d-acce-01f94976dac1
base	
spark-mllib_2.3	2e51f700-bca0-4b0d-88dc-5c6791338875
base	
pytorch-onnx_1.1-py3.6-edt	32983cea-3f32-4400-8965-dde874a8d67e
base	
spark-mllib_3.0-py37	36507ebe-8770-55ba-ab2a-eafe787600e9
base	

spark-mllib_2.4	390d21f8-e58b-4fac-9c55-d7ceda621326
base	
autoai-ts_rt22.2-py3.10	396b2e83-0953-5b86-9a55-7ce1628a406f
base	
xgboost_0.82-py3.6	39e31acd-5f30-41dc-ae44-60233c80306e
base	
pytorch-onnx_1.2-py3.6-edt	40589d0e-7019-4e28-8daa-fb03b6f4fe12
base	
pytorch-onnx_rt22.2-py3.10	40e73f55-783a-5535-b3fa-0c8b94291431
base	
default_r36py38	41c247d3-45f8-5a71-b065-8580229facf0
base	
autoai-ts_rt22.1-py3.9	4269d26e-07ba-5d40-8f66-2d495b0c71f7
base	
autoai-obm_3.0	42b92e18-d9ab-567f-988a-4240ba1ed5f7
base	
pmml-3.0_4.3	493bcb95-16f1-5bc5-bee8-81b8af80e9c7
base	
spark-mllib_2.4-r_3.6	49403dff-92e9-4c87-a3d7-a42d0021c095
base	
xgboost_0.90-py3.6	4ff8d6c2-1343-4c18-85e1-689c965304d3
base	
pytorch-onnx_1.1-py3.6	50f95b2a-bc16-43bb-bc94-b0bed208c60b
base	
autoai-ts_3.9-py3.8	52c57136-80fa-572e-8728-a5e7cbb42cde
base	
spark-mllib_2.4-scala_2.11	55a70f99-7320-4be5-9fb9-9edb5a443af5
base	
spark-mllib_3.0	5c1b0ca2-4977-5c2e-9439-ffd44ea8ffe9
base	
autoai-obm_2.0	5c2e37fa-80b8-5e77-840f-d912469614ee
base	
spss-modeler_18.1	5c3cad7e-507f-4b2a-a9a3-ab53a21dee8b
base	
cuda-py3.8	5d3232bf-c86b-5df4-a2cd-7bb870a1cd4e
base	
runtime-22.2-py3.10-xc	5e8cddff-db4a-5a6a-b8aa-2d4af9864dab
base	
autoai-kb_3.1-py3.7	632d4b22-10aa-5180-88f0-f52dfb6444d7
base	

Note: Only first 50 records were displayed. To display more use
 'limit' parameter.

```
software_spec_uid =
client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid
```

'12b83a17-24d8-5082-900f-0ab31fbfd3cb'

